

The CSM Agent — Proactive Agentic Adoption Journey

A collection of AI agents that monitors product usage, identifies where customers are not getting full value, and delivers personalised guidance at scale — with CSMs reviewing the interactions that need their judgment.

Executive Summary What we're building and the plan	The Story Why it matters, for whom	Engineering Detail How it works, component by component	ExCo Submission Appendix — Big Bet template
--	--	---	---

Engineering Overview

Eight layers. **Data platform: Snowflake** — signals are computed here, and Snowflake Cortex (an intelligent search capability built into Snowflake) enables agents to search knowledge bases by meaning rather than by keyword. **Snowflake Cortex is not the agent-building tool** — it is the search service the agents call when they need to find relevant information. **Agent coordination: Microsoft Copilot Studio** — all five agents are built and coordinated here. **Delivery: Gainsight CS and Gainsight PX** — native connections live April 2026. World Check 1's completed Gainsight integration is the reference implementation.

Snowflake Cortex vs Microsoft Copilot Studio — understanding the distinction: Copilot Studio is where the agents are built. It decides what to do and when. Snowflake Cortex is a search capability within Snowflake that agents query when they need to find relevant information — a customer's feedback history, the most relevant approved content, a CSM's past messages. Think of Copilot Studio as the decision-maker and Cortex as the library it searches. Neither replaces the other.

Expand / Collapse All

1 Raw Data Sources

FOUNDATION

Where the data comes from — usage events, customer feedback, account context, and product releases

STRUCTURED

Product Usage Events

Login timestamps, feature interaction events, workflow completions, session lengths. Event-level detail. Already in Snowflake.

Refresh: daily batch or near-real-time for time-sensitive signals, using Snowflake Dynamic Tables.

UNSTRUCTURED

VOC Golden Source

Qualtrics survey text, Salesforce incidents and enhancement requests, Gainsight CSM health notes, Asana escalations, Salesloft call transcripts. All aggregated in Snowflake. Tagging work already started via the QAS mechanism; deeper build-out partially in flight.

Confirm scope and timeline for the VOC tagging build-out with the Corporate Engineering AI team. The key question is whether planned tagging covers the full

BEHAVIOURAL

Gainsight PX Engagement History

What in-product content each user has been shown, clicked, or dismissed. Prevents the system repeating content a customer has already seen. Available via Gainsight PX REST API (production-ready).

Pull into Snowflake daily. Use PX connection for real-time queries at agent decision time.

CRM

Account and Relationship Data

Renewal dates, contract value, account tier, influence score, health scores, success plan status. Lives in Gainsight CS and Salesforce.

Queried live at decision time via Gainsight CS connection. Salesforce data via Corporate Engineering AI team's planned Salesforce connection (Q2 — alignment required).

PRODUCT

Tagged Enhancement Releases

Every product release tagged across six fields: module, persona, workflow impacted, problem solved, behaviour trigger (what usage pattern indicates a user would benefit), and adoption complexity. Requires a new step in the product release process.

Stored in Snowflake. The behaviour trigger field is critical — it is how the system identifies which users a given enhancement is relevant to.

breadth: surveys, incidents, health notes, and transcripts.

CSM

CSM Decision Log

A record of every accept, edit, or discard from the CSM review queue. Edits stored alongside the original so the system can identify what changed and why. Feeds voice archive improvement and routing rule refinement over time.

The difference between original and edited versions is stored. A nightly process classifies the type of change: tone, content, length, or factual correction.

Design principle: Pre-compute and store as much as possible in Snowflake before agents query it. This reduces wait time and failure risk. Exception: Gainsight CS account context — always read live to ensure accuracy.



2

Signal Engineering

Making sense of the data — turning raw usage events into meaningful signals that tell the agents when and why to act

MUST BUILD · SNOWFLAKE

Built as scheduled SQL jobs in Snowflake, organised using dbt (an open-source tool that structures SQL logic into testable, version-controlled models — the standard approach for this type of transformation work). All output lands in a single signals table with a consistent structure. Every agent queries one place, regardless of signal type.

COMPUTE

Usage Metric Tables

Per-user metrics refreshed daily: days since last login, percentage of core features used, workflows completed vs available, content consumed, onboarding milestone completion rate.

Snowflake scheduled tasks or dbt models. Dynamic Tables for signals needing more frequent refresh.

LOGIC

Behavioural Pattern Detection

SQL rules defined by the Working Group: for example, login frequency down more than 30% over three weeks, AND a core workflow never completed, AND renewal within 180 days. Each rule that fires produces a row in the signals table.

Start rule-based in the pilot. Once six months of outcome data exists, consider a predictive model using

SCORING

Signal Severity Scoring

Each signal scored 1–10, calibrated against domain expertise initially and refined against outcome data as the pilot generates it. Green (no action needed), Amber (nudge), Red (intervene).

Rule-based weights for the pilot. Post-pilot: a predictive model using Snowflake's built-in machine learning capabilities.

ENRICHMENT

Persona and Journey Stage Tagging

Each signal enriched at the point of computing: user role, account tier, journey stage (onboarding, active, at-risk, approaching renewal), and influence score. Joined from CRM context.

Start rule-based using job title normalisation and bucketing. Evolve to pattern-based grouping after the pilot generates sufficient data.

MATCHING

Enhancement Relevance Matching

Compares tagged enhancement releases against user behaviour profiles. For each new release, identifies users whose usage matches the enhancement's behaviour trigger field. Produces a relevance score per user per enhancement. Only fires for users who have not already been shown that enhancement in-product.

Snowflake's built-in machine learning capabilities — no external tooling needed.

Scored by match strength relative to adoption complexity. Prevents sending a message about a feature a user has already engaged with.

Key design principle: All signal types land in one signals table with a consistent structure. Downstream agents query one place. Normalise early — inconsistent structures across signal types create significant complexity later in the build.



3

Searchable Knowledge Bases
Giving agents the right context before they act — four stores searched by meaning, not by keyword, using Snowflake Cortex

MUST BUILD · SNOWFLAKE CORTX

▲

This is where Snowflake Cortex is used. Cortex provides intelligent search built into Snowflake — agents can search large bodies of text by meaning rather than exact keywords. It is not used to build or run agents; it is a search service agents call when they need to find relevant information. All four knowledge bases below live inside Snowflake, meaning no additional database or search tooling is required.

The four knowledge bases — what they contain and why intelligent search is needed for each

- **Customer feedback and history (VOC data)** — survey comments, call summaries, CSM health notes. The agent searches this to understand what a specific customer has said and felt. Keyword matching would miss too much. Tagging via QAS mechanism already started; scope of planned build-out to be confirmed with the Corporate Engineering AI team.
- **Approved content and playbooks** — email templates, playbook summaries, enhancement descriptions. The agent finds the most relevant approved content for this signal and user type. Relevance is contextual, not a simple lookup.
- **Gainsight PX engagement history** — what content a user has already been shown in-product. The agent checks this to avoid repetition and identify what has worked before.
- **CSM voice archive** — accepted and unedited messages per CSM, used as style anchors when drafting. The agent retrieves examples matching the current message type and user to write in the CSM's voice. Grows with every accepted draft from the review queue.

SEARCH 1

LEVERAGE

Customer Feedback — VOC Data

Customer feedback, call summaries, and CSM health notes made searchable by meaning. VOC data tagging via the QAS mechanism already started and partially in flight. This builds on existing work rather than starting from scratch.

Confirm current and planned tagging scope with the Corporate Engineering AI team. The key question is whether it covers the full breadth of the VOC golden source across all source types.

SEARCH 2

Approved Content and Playbooks

Approved templates, playbooks, and enhancement descriptions made searchable. Agent retrieves the most relevant approved content for the current signal and user type. Metadata filters — role, channel, adoption complexity — applied before similarity ranking.

Must be built in Snowflake Cortex. Add a performance score field, updated weekly from outcome

SEARCH 3

PX Engagement History

Descriptions of in-product content previously shown per user, made searchable. Agent checks this before recommending in-product content to avoid repeating what a user has already seen or dismissed.

Lightweight. Short descriptions. Scoped by user. Refreshed daily from PX REST API.

SEARCH 4

CSM Voice Archive

Accepted and unedited messages per CSM, made searchable. Agent retrieves style examples matching the current message type and user as drafting anchors. Starts empty and grows with every accepted draft. Approximately 20–30 accepted messages per CSM before voice matching becomes reliable.

Scoped by CSM. Start collecting from pilot day one — the earlier this starts, the better the voice quality at full rollout.

data, to weight retrieval toward content that has proven effective.

All four knowledge bases live inside Snowflake. Snowflake Cortex provides the intelligent search capability natively. No separate search tooling is needed — a meaningful reduction in architectural complexity and ongoing maintenance.



4 **Signals and Content Library**

The decision rulebook — what each signal means, what to do about it, and what content the agents draw from

MUST BUILD — HUMAN + ENGINEERING

RULES TABLE

Signal to Action Mapping

A Snowflake table that maps each signal type, user type, and severity level to an approved response: what kind of message, which channel, which content to draw from, and whether a CSM needs to review before sending. The Working Group updates this table directly — no code change required.

Separating this logic into a data table (rather than embedding it in agent code) is an important architectural decision. It means the Working Group can adjust what the system does as the pilot generates evidence, without engineering involvement.

CONTENT

Enhancement Content Blocks

A short, role-appropriate description of each product enhancement, authored by the content organisation. Stored in Snowflake and made searchable via Knowledge Base 2. Agents retrieve and personalise these — they never generate product descriptions from scratch, which prevents inaccurate claims reaching customers.

Target: authored within 48 hours of each product release. Versioned — old versions archived, not deleted, so the system can understand what was communicated when.

TEMPLATES

Email and In-Product Templates

Structural message templates per message type, stored in Gainsight Journey Orchestrator. Variable fields — customer name, feature name, workflow context — are populated by the agent before triggering the send via the Gainsight connection. Agents never write directly into Gainsight.

Confirm Gainsight Journey Orchestrator supports fully dynamic body content, not just standard name and date fields. Verify this before building templates.

RULES TABLE

CSM Review Routing Rules

A separate Snowflake table defining the conditions under which a message goes to a CSM for review rather than sending automatically: high severity, high influence, frustration detected, complex adoption topic, or strategic account tier. The Working Group adjusts these thresholds without any code change.

Stored in a dedicated rules table — not hardcoded into agent logic. Thresholds should be reviewed monthly during the pilot as outcome data accumulates.

Architecture principle: Rules and routing logic live in Snowflake tables, not embedded in the agents. The Working Group controls what the system does without engineering involvement. Keep agent logic minimal; keep decisions in data.



5 **The Five Agents**

Each agent has one clearly defined job — they work in sequence, with defined handoffs between them, all built and coordinated in Microsoft Copilot Studio

MUST BUILD · COPILOT STUDIO

AGENT 1

Signal Detection Agent

AGENT 2

VOC Personalisation Agent

AGENT 3

Next Best Action Agent

AGENT 4

EARLY WIN

Content Build Agent

AGENT 5

EARLY WIN

Monitors the signals table on a defined schedule. Filters for signals above the severity threshold and passes the signal details, user, account, and severity level to the next agent. No AI generation involved — pure detection logic.

Copilot Studio autonomous trigger connected to Snowflake. Uses the Snowflake Query capability registered in Skills Middleware.

Gathers everything needed before a decision is made: usage signals from Snowflake, customer feedback via the VOC knowledge base, engagement history via the PX knowledge base, and live account data via the Gainsight connection. Produces a structured context summary passed to the next agent.

Parallel data retrieval where possible to reduce wait time. AI generation used only to summarise VOC search results — not for decisions.

Queries the signal-to-action mapping table and routing rules table. Returns the message type, delivery channel, content to draw from, and whether a CSM review is required. This is a pure rules lookup — no AI generation. Decisions are consistent and fully auditable.

Deterministic logic only. If routing logic grows complex, a decision tree is preferred over an AI prompt. Auditability is non-negotiable for a customer-facing system.

Retrieves relevant content blocks from the approved content knowledge base and voice examples from the CSM voice archive, then generates a personalised draft. Includes a quality and assurance check against the original request before routing — ensuring the message meets the intent of whoever submitted it (for example, a product manager submitting an enhancement for customer communication). Constrained to retrieved content — cannot generate product claims from scratch.

MVP version to be delivered early for use in current digital CX promotional campaigns through the joint intake and governance process. Evolves over time as the voice archive and content library grow. AI model TBD — to be aligned with the Corporate Engineering AI team.

Assessment and Prioritization Agent
Handles two related jobs: assessing and prioritising incoming requests from the product organisation for customer communications (ensuring the most impactful messages go first, based on customer impact, problem severity, business value, and execution feasibility); and routing outbound drafted messages to the appropriate CSM for review where conditions require it.

Early version prioritises the manual joint intake and governance process already in place. Evolves toward a fully AI-driven intake as volume grows and the system matures. Queue ordered by urgency and confidence.

Agent capabilities required — to be registered in the Corporate Engineering AI team's Skills Middleware

Snowflake Query

Read from the signals table, action mapping, routing rules, and content library. Used by Agents 1, 2, and 3. Requirements to be aligned with the Corporate Engineering AI team's Snowflake API design.

Gainsight CS Connection

Read account context and health data. Create review tasks in the CSM inbox. Trigger email sends via Journey Orchestrator. Used by Agents 2, 5, and the delivery layer. Coverage to be confirmed via the MCP Gateway.

AI Draft Generation

Generate personalised message drafts from structured inputs, constrained to retrieved content. Agent 4 only. AI model to be confirmed — alignment with Corporate Engineering AI team recommended to avoid duplicating infrastructure.

Intelligent Search (Snowflake Cortex)

Search across Knowledge Bases 2, 3, and 4 by meaning. Customer feedback search (Knowledge Base 1) via QAS if the Corporate Engineering AI team's tagging build-out covers the full VOC golden source.

Gainsight PX Connection

Trigger in-product messages for specific users. Read engagement history. Used by the delivery layer and Agent 2. Coverage to be confirmed via the MCP Gateway.

Snowflake Write

Write outcome data and CSM decisions back to Snowflake for the learning loop. Used by the delivery layer and the CSM review interface after each action.

Note on agent coordination: Copilot Studio's multi-agent coordination capability supports this pattern — a parent orchestrator routes to specialised agents in sequence. This capability is relatively new and the handoff patterns between agents will require careful build and testing. Budget adequate time for this phase.

UX

Prioritised Daily Inbox

A queue of drafts requiring CSM review, sorted by urgency and confidence. Each item shows the customer, a brief explanation of what triggered it, the message type, a confidence indicator, and the full draft inline. Designed to be scanned quickly.

Start with a Gainsight dashboard section in the pilot. Invest in a richer inline editing interface once CSM usage patterns are better understood. Do not over-engineer before knowing how CSMs want to interact with it.

ACTIONS

Accept / Edit / Discard

Accept: triggers delivery immediately. Edit: inline changes, with the original preserved alongside the final version. Discard: archived with an optional reason selected from a short dropdown. All three actions written to the decision log in Snowflake.

The difference between original and edited versions is stored. Nightly process classifies the type of change. Factual corrections trigger a content library review within 24 hours.

LOGIC

Confidence Indicator

Each draft shows a confidence band — High, Medium, or Review Carefully — based on how many signals converged, the quality of context matches, and the historical acceptance rate for this type of message. Context for the CSM, not a directive.

"Review Carefully" signals agent uncertainty and draws more prominent attention. Raw scores stored internally for calibration purposes.

GOVERNANCE

Automation Controls

A Working Group admin panel for adjusting which message types require CSM review versus sending automatically. When a message type's unedited acceptance rate exceeds a configurable threshold over 30 days, the Working Group can move it to fully automated. Full audit trail of all changes.

Thresholds stored in the routing rules table in Snowflake, not in code. Changing a threshold is a data update, not a code deployment.

If reviewing a message takes more than a minute, CSMs will stop using the queue. Speed of the interface is a product quality requirement. Test with a small group of CSMs before adding features or scaling.

CHANNEL 1

Email via Journey Orchestrator

Agents trigger a Journey Orchestrator programme via the Gainsight CS connection, passing personalised content variables. Sent from the CSM's email identity. Works for both automated sends and sends that follow CSM review.

CHANNEL 2

In-Product via Gainsight PX

Agents trigger a targeted in-product prompt, tooltip, or guide for a specific user, appearing at their next login. Uses PX's existing engagement engine — no custom front-end work needed where PX is already instrumented. World Check 1 already integrated;

CHANNEL 3

CSM-Sent Following Review

When a CSM accepts or edits a draft in the review queue, the system triggers the Journey Orchestrator send using the final version. The customer receives it as a personal email from their CSM — because it was reviewed and approved by them.

CHANNEL 4

Prepared Brief for CSM Outreach

When escalation conditions are met — a high-influence, frustrated customer with a complex topic — a task is created in Gainsight CS for the responsible CSM, pre-loaded with signal context, a conversation guide drawn from the content library, and a

Confirm Gainsight Journey Orchestrator supports fully dynamic body content, not just standard name and date fields. Verify before building templates.

Workspace integration in progress.

Gainsight PX REST API: trigger engagement by user ID. World Check 1's completed integration is the reference for the Workspace build.

Same connection as Channel 1. Trigger source marked to distinguish CSM-approved sends from automated ones. The final version — not the original draft — is always what sends.

suggested outcome. The CSM leads the conversation; the agents handled the preparation.

Gainsight CS connection: create a task with account, assignee CSM, due date, and context. Conversation guide retrieved from content library search and attached as task detail.

Both Gainsight connections are live as of April 2, 2026. Confirm connection endpoint coverage against the specific actions above before scoping any custom development. World Check 1's completed integration is the reference — use it to validate capability before the Workspace build begins.



System Learning and Improvement

8

The whole CSM Agent needs to keep improving. This layer feeds outcomes and decisions back into every component — signals, content, agent behaviour, and routing rules — so the system gets more precise over time

CONTINUOUS · ALL LAYERS

This layer is distinct from the five agents above. It does not sit in sequence — it runs continuously alongside the whole system, feeding outcomes and decisions back into each of the layers above independently. Every agent has its own improvement mechanism; this layer is what ensures those improvements happen and that the system as a whole is moving in the right direction for customers.

Customer Outcome Tracking

After each message: record whether the customer opened it, clicked through, and adopted the relevant feature within 14 and 30 days. This is the primary measure of whether the system is working for customers.

Feeds back into: Signal Engineering (Layer 2) — calibrates which signals actually predict adoption improvement. Content Library (Layer 4) — surfaces which content is driving outcomes.

CSM Decision Tracking

Every accept, edit, and discard from the CSM review queue is logged. Edits classified by type nightly. Factual corrections trigger a content library review within 24 hours. High discard rates for a signal type trigger a Working Group review.

Feeds back into: CSM Voice Archive (Layer 3, Search 4) — accepted unedited messages improve voice matching. Routing Rules (Layer 4) — high discard rates indicate routing thresholds need adjustment.

Signal Threshold Calibration

Monthly: compare signals against adoption outcomes. For each signal type, calculate how much adoption improved following an intervention, and what the CSM acceptance rate was. Low-performing signals are reviewed and adjusted. High-performing ones become candidates to send automatically.

Feeds back into: Signal Engineering (Layer 2) — threshold and rule refinement. Signal to Action Mapping (Layer 4) — routing decisions updated based on evidence.

Content Performance Scoring

Per content block: track sends, opens, click-throughs, 30-day adoption rate, and CSM acceptance rate. High-performing content is weighted higher in search retrieval. Low-performing content is flagged for the content organisation to review or retire.

Feeds back into: Approved Content Knowledge Base (Layer 3, Search 2) — performance score added to search weighting. Enhancement Content Blocks (Layer 4) — poor performers flagged for update.

Voice Archive Growth

Accepted, unedited messages added to each CSM's voice archive nightly. Track how often CSMs edit drafts over time as the key quality metric — this should decrease as

Working Group Monthly Review

The CSAM AI Working Group reviews performance data monthly: which signals are producing adoption improvement, which content is landing, what CSM edit rates look

the archive grows and the agent learns ech CSM's style. No model retraining needed; the search-based approach improves naturally.

Feeds back into: CSM Voice Archive (Layer 3, Search 4) — growing archive improves voice matching quality for the Content Build Agent (Layer 5).

like, and which message types are candidates for automation graduation. Human oversight of the learning loop is non-negotiable.

Feeds back into: All layers — routing rules, signal thresholds, content library, and automation controls are all adjusted based on Working Group decisions.

Both feedback cycles must be reported to the Working Group monthly from day one of the pilot. High adoption impact combined with low CSM acceptance rates means the system cannot scale — CSMs will bypass the review queue. Both need to be healthy simultaneously, and the Working Group is the governance layer that ensures this.

✓ **Already Have or Can Use**

- Usage data and VOC data, both in Snowflake
- VOC tagging via QAS mechanism already started; deeper build-out partially in flight
- Skills Middleware — Corporate Engineering AI team (completed Q1 2026)
- Gainsight CS connection (live April 2026)
- Gainsight PX connection and REST API (live)
- World Check 1 as a reference for the complete Gainsight PX and CS integration
- Copilot Studio with multi-agent coordination capability
- Snowflake Cortex — intelligent search capability built into Snowflake
- Gainsight Journey Orchestrator for email delivery

✗ **Must Build**

- Signal detection rules and scoring logic in Snowflake
- Enhancement relevance matching logic
- Three knowledge bases in Snowflake Cortex: approved content, PX history, CSM voice archive
- Signal-to-action mapping and routing rules tables
- Enhancement content block authoring workflow (content org)
- Content Build Agent — MVP early, full version with voice matching later
- Assessment and Prioritization Agent — MVP for intake prioritization, full version over time
- Signal Detection, VOC Personalisation, and Next Best Action agents in Copilot Studio
- CSM review queue interface in Gainsight
- CSM decision log and nightly edit classification
- Outcome tracking and system learning infrastructure

Synergies with LSEG Assist — Intelligent Enterprise

The Corporate Engineering AI team is building infrastructure that directly overlaps with what the CSM Agent requires. Using this shared infrastructure reduces build scope, avoids duplication, and accelerates delivery.

The LSEG Assist roadmap includes a strand called "**Intelligent Enterprise: Agent ready data and tools.**" Its purpose is to make LSEG's data and systems accessible to any agent via a common set of connections (MCP — a standard protocol that lets AI agents connect to external systems and data sources) and reusable agent capabilities (Skills). This is the infrastructure layer the CSM Agent sits on top of. Four components are directly relevant.

LIVE — IN PROGRESS

QAS — Enterprise Search as a Service

The Corporate Engineering AI team runs a managed intelligent search service called QAS, indexing 37+ datasets with 2,300 users per week. VOC data tagging

COMPLETED Q1 2026

Skills Middleware

A completed layer that makes agent capabilities available in a standard format across LSEG systems. The CSM Agent requires six specific capabilities — querying

using the QAS mechanism is already started and a deeper build-out is partially in flight.

Replaces one of the four planned knowledge base builds. Remaining work: confirm the planned tagging build-out covers the full VOC golden source breadth needed for the CSM Agent — surveys, incidents, health notes, and transcripts.

Snowflake, intelligent search, Gainsight CS, Gainsight PX, AI draft generation, and writing back to Snowflake.

CSM Agent capabilities can be registered in this shared framework rather than built and maintained independently. Standard format, shared governance, reduced duplication.

Q2 2026 — IN PROGRESS

Snowflake Data APIs

The Corporate Engineering AI team is building Snowflake connectivity via their standard connection framework. The CSM Agent's primary data needs — signals, customer feedback, and the content library — all require exactly this.

Early alignment is needed. CSM Agent data requirements — the specific tables and access patterns needed — should inform the Snowflake API design before it is finalised. Retrofitting after the fact is costly.

Q2 2026 — NOT YET STARTED

Salesforce Native Connection

The Corporate Engineering AI team is planning a Salesforce connection to surface account details and pipeline data. The CSM Agent requires exactly this for the VOC Personalisation Agent — CRM account data, renewal dates, and influence scores.

Early alignment with the Corporate Engineering AI team is needed. The CSM Agent use case and the LSEG Assist use case both require this connection — shared delivery is more efficient than parallel builds.

Q3 2026 — PLANNED

Dynamic Snowflake Queries

A planned capability that executes a single efficient query against Snowflake rather than multiple sequential ones, reducing cost and processing time. The Signal Detection Agent makes repeated Snowflake queries.

Not a dependency for launch, but a meaningful performance improvement once live. Signal query design should account for this from the start.

CONFIRMATION REQUIRED

MCP Gateway — Enterprise Security Layer

The Corporate Engineering AI team is building a secure gateway that handles authentication and governance for all agent-to-system connections. The CSM Agent requires connections to Gainsight CS and Gainsight PX — both vendor-hosted tools requiring enterprise-grade security controls.

Confirmation needed: are Gainsight CS and PX in scope for the MCP Gateway? If not, a clear approved security path for those connections must be agreed before build begins.

Specific Alignment Points with the Corporate Engineering AI Team

- 1 Confirm VOC tagging scope and timeline.** Tagging via the QAS mechanism is already started. Confirmation needed on whether the planned build-out covers the full VOC golden source — survey text, Salesforce incidents and enhancement requests, Gainsight health notes, and Salesloft call transcripts — and when this will be complete. **In Progress — confirm scope**
- 2 Include CSM Agent data requirements in the Snowflake API design.** The CSM Agent requires read access to three specific data areas: the signals table, the content library, and the VOC golden source. These requirements should be incorporated into the Snowflake API design while it is still in progress, to avoid retrofitting later. **Q2 — in progress**
- 3 Align on the Salesforce connection build.** Both the CSM Agent and LSEG Assist require this connection. Early alignment between teams avoids two parallel builds for the same capability. The connection is not yet started on the Corporate Engineering AI team's Q2 roadmap. **Q2 — not yet started**

4 **Confirm Gainsight CS and PX coverage in the MCP Gateway.** Both connections are used extensively by the CSM Agent delivery layer. Whether they are covered by the gateway, or what the approved alternative security path is, needs to be confirmed before the delivery layer build begins. [Confirm now](#)

5 **Register CSM Agent capabilities in Skills Middleware.** Skills Middleware is complete. The six capabilities the CSM Agent requires — Snowflake query, intelligent search, Gainsight CS, Gainsight PX, AI draft generation, and Snowflake write — should be defined and registered in the shared framework rather than built and maintained independently. [Ready to start](#)